


Autoregressive → block diffusion training


Qwen3.8-27B · 16× H200

Training time per 1B tokens



Block diffusion fine-tuning


DiffusionGemma 26B-A4B · 16× H200



Context length

Speculative decoding drafter training

DFlash2 for Qwen3.8-27B · 8× H100



● Context-Shared Block Parallelism (CSBP)

○ Best existing parallelism strategies