






(a) The **Idmatrix** instruction ensures that each thread gets what it needs for compute in W8A8 GEMM



(b) However, the **Idmatrix** instruction fails for W4A8 GEMM due to **storage-compute mismatch**



(c) Our solution: **compute-aware weight reorder**