

Rollout engine — SGLang

dedicated GPU pool, batch-wise sampling

rollout batch 1
sampled with θ_0 rollout batch 2
sampled with θ_0 rollout batch 3
sampled with θ_8 rollout batch 4
sampled with θ_{16} **Training engine — Megatron**

dedicated GPU pool, batch-wise training

 θ_0 θ_8 θ_{16} θ_{24}

8 optimizer steps

8 optimizer steps

8 optimizer steps

→ rollout data flow

→ update_weights() broadcast

→ optimizer steps

weight-broadcast instant