


one shared GPU pool — the two engines take turns, never overlapping



staleness $n = 0$ — every batch is trained by the exact weight that sampled it;
the price: rollout, training and weight broadcast run strictly in serial.

→ rollout data flow → update_weights() broadcast □ idle engine (serial cost)